

DAY 10:

Bottle Dataset

In [1]:

```
#to import libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

In [2]:

```
df=pd.read_csv(r"E:\Dataset\9_bottle.csv")[0:500]  
df
```

```
C:\ProgramData\Anaconda3\lib\site-packages\IPython\core\interactiveshell.p  
y:3165: DtypeWarning: Columns (47,73) have mixed types.Specify dtype optio  
n on import or set low_memory=False.  
    has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```

Out[2]:

| Cst_Cnt | Btl_Cnt | Sta_ID | Depth_ID | Depthm | T_degC | Salnty | O2ml_L | STheta | O2Sat |
|---------|---------|--------|----------------------------------------------------------|--------|--------|--------|--------|--------|-------|
| 0 | 1 | 1 | 054.0 056.0 19-4903CR-HY-060-0930-05400560-0000A-3 | 0 | 10.50 | 33.440 | NaN | 25.649 | NaN |
| 1 | 1 | 2 | 054.0 056.0 19-4903CR-HY-060-0930-05400560-0008A-3 | 8 | 10.46 | 33.440 | NaN | 25.656 | NaN |
| 2 | 1 | 3 | 054.0 056.0 19-4903CR-HY-060-0930-05400560-0010A-7 | 10 | 10.46 | 33.437 | NaN | 25.654 | NaN |
| 3 | 1 | 4 | 054.0 056.0 19-4903CR-HY-060-0930-05400560-0019A-3 | 19 | 10.45 | 33.420 | NaN | 25.643 | NaN |
| 4 | 1 | 5 | 054.0 056.0 19-4903CR-HY-060-0930-05400560-0020A-7 | 20 | 10.45 | 33.421 | NaN | 25.643 | NaN |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 495 | 16 | 496 | 063.3 058.0 19-4903CR-HY-065-1030-06330580-0700A-7 | 700 | 4.90 | 34.269 | NaN | 27.114 | NaN |
| 496 | 16 | 497 | 063.3 058.0 19-4903CR-HY-065-1030-06330580-0792A-3 | 792 | 4.50 | 34.310 | NaN | 27.191 | NaN |
| 497 | 16 | 498 | 063.3 058.0 19-4903CR-HY-065-1030-06330580-0800A-7 | 800 | 4.48 | 34.311 | NaN | 27.194 | NaN |
| 498 | 16 | 499 | 063.3 058.0 19-4903CR-HY-065-1030-06330580-0900A-7 | 900 | 4.21 | 34.319 | NaN | 27.230 | NaN |

In [31]:

```
Cst_Cnt  Btl_Cnt  Sta_ID  Depth_ID  Depthm  T_degC  Salnty  O2ml_L  STheta  O2Sat
df.info()
```

| | | | | | | | | | | |
|-----|----|-----|----------------|----------------------------------------|------|------|--------|-----|--------|-----|
| | | | | 19-4903CR-HY-065-1030-06330580-1000A-7 | | | | | | |
| 499 | 16 | 500 | 063.3 058.0 | | 1000 | 3.95 | 34.329 | NaN | 27.265 | NaN |

500 rows × 74 columns

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 500 entries, 0 to 499
```

```
Data columns (total 74 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|-------------|----------------|---------|
| 0 | Cst_Cnt | 500 non-null | int64 |
| 1 | Btl_Cnt | 500 non-null | int64 |
| 2 | Sta_ID | 500 non-null | object |
| 3 | Depth_ID | 500 non-null | object |
| 4 | Depthm | 500 non-null | int64 |
| 5 | T_degC | 499 non-null | float64 |
| 6 | Salnty | 494 non-null | float64 |
| 7 | O2ml_L | 0 non-null | float64 |
| 8 | STheta | 493 non-null | float64 |
| 9 | O2Sat | 0 non-null | float64 |
| 10 | Oxy_μmol/Kg | 0 non-null | float64 |
| 11 | BtlNum | 0 non-null | float64 |
| 12 | RecInd | 500 non-null | int64 |
| 13 | T_prec | 499 non-null | float64 |
| 14 | T_qual | 4 non-null | float64 |
| 15 | S_prec | 494 non-null | float64 |
| 16 | S_qual | 10 non-null | float64 |
| 17 | P_qual | 500 non-null | float64 |
| 18 | O_qual | 500 non-null | float64 |
| 19 | SThtaq | 14 non-null | float64 |
| 20 | O2Satq | 500 non-null | float64 |
| 21 | ChlorA | 0 non-null | float64 |
| 22 | Chlqua | 500 non-null | float64 |
| 23 | Phaeop | 0 non-null | float64 |
| 24 | Phaqua | 500 non-null | float64 |
| 25 | PO4uM | 0 non-null | float64 |
| 26 | PO4q | 500 non-null | float64 |
| 27 | SiO3uM | 0 non-null | float64 |
| 28 | SiO3qu | 500 non-null | float64 |
| 29 | NO2uM | 0 non-null | float64 |
| 30 | NO2q | 500 non-null | float64 |
| 31 | NO3uM | 0 non-null | float64 |
| 32 | NO3q | 500 non-null | float64 |
| 33 | NH3uM | 0 non-null | float64 |
| 34 | NH3q | 500 non-null | float64 |
| 35 | C14As1 | 0 non-null | float64 |
| 36 | C14A1p | 0 non-null | float64 |
| 37 | C14A1q | 500 non-null | float64 |
| 38 | C14As2 | 0 non-null | float64 |
| 39 | C14A2p | 0 non-null | float64 |
| 40 | C14A2q | 500 non-null | float64 |
| 41 | DarkAs | 0 non-null | float64 |
| 42 | DarkAp | 0 non-null | float64 |
| 43 | DarkAq | 500 non-null | float64 |
| 44 | MeanAs | 0 non-null | float64 |
| 45 | MeanAp | 0 non-null | float64 |
| 46 | MeanAq | 500 non-null | float64 |
| 47 | IncTim | 0 non-null | object |
| 48 | LightP | 0 non-null | float64 |
| 49 | R_Depth | 500 non-null | float64 |
| 50 | R_TEMP | 499 non-null | float64 |
| 51 | R_POTEMP | 495 non-null | float64 |
| 52 | R_SALINITY | 494 non-null | float64 |
| 53 | R_SIGMA | 486 non-null | float64 |
| 54 | R_SVA | 486 non-null | float64 |
| 55 | R_DYNHT | 500 non-null | float64 |

```

56 R_O2                0 non-null    float64
57 R_O2Sat             0 non-null    float64
58 R_SIO3              0 non-null    float64
59 R_PO4               0 non-null    float64
60 R_NO3               0 non-null    float64
61 R_NO2               0 non-null    float64
62 R_NH4               0 non-null    float64
63 R_CHLA              0 non-null    float64
64 R_PHAEO             0 non-null    float64
65 R_PRES              500 non-null  int64
66 R_SAMP              0 non-null    float64
67 DIC1                0 non-null    float64
68 DIC2                0 non-null    float64
69 TA1                 0 non-null    float64
70 TA2                 0 non-null    float64
71 pH2                 0 non-null    float64
72 pH1                 0 non-null    float64
73 DIC Quality Comment 0 non-null    object

```

dtypes: float64(65), int64(5), object(4)

memory usage: 289.2+ KB

In [4]:

```
df.columns
```

Out[4]:

```

Index(['Cst_Cnt', 'Btl_Cnt', 'Sta_ID', 'Depth_ID', 'Depthm', 'T_degC',
      'Salnty', 'O2ml_L', 'STheta', 'O2Sat', 'Oxy_μmol/Kg', 'BtlNum',
      'RecInd', 'T_prec', 'T_qual', 'S_prec', 'S_qual', 'P_qual', 'O_qual',
      'SThta', 'O2Satq', 'ChlorA', 'Chlqua', 'Phaeop', 'Phaqua', 'P04u',
      'P04q', 'SiO3uM', 'SiO3qu', 'NO2uM', 'NO2q', 'NO3uM', 'NO3q', 'NH3u',
      'NH3q', 'C14As1', 'C14A1p', 'C14A1q', 'C14As2', 'C14A2p', 'C14A2q',
      'DarkAs', 'DarkAp', 'DarkAq', 'MeanAs', 'MeanAp', 'MeanAq', 'IncTi',
      'LightP', 'R_Depth', 'R_TEMP', 'R_POTEMP', 'R_SALINITY', 'R_SIGMA',
      'R_SVA', 'R_DYNHT', 'R_O2', 'R_O2Sat', 'R_SIO3', 'R_PO4', 'R_NO3',
      'R_NO2', 'R_NH4', 'R_CHLA', 'R_PHAEO', 'R_PRES', 'R_SAMP', 'DIC1',
      'DIC2', 'TA1', 'TA2', 'pH2', 'pH1', 'DIC Quality Comment'],
      dtype='object')

```

Linear Regression

In [5]:

```

x=df[['Cst_Cnt', 'Btl_Cnt', 'Depthm', 'RecInd', 'P_qual', 'O_qual', 'O2Satq']]
y=df['R_PRES']

```

In [6]:

```
# to split my dataset into test and train data
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

In [7]:

```
from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[7]:

LinearRegression()

In [8]:

```
print(lr.score(x_test,y_test))
```

0.9999987632216758

In [9]:

```
lr.score(x_train,y_train)
```

Out[9]:

0.9999989709800864

Ridge Regression

In [10]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [11]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

Out[11]:

0.9999987078970033

Lasso Regression

In [12]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[12]:

Lasso(alpha=10)

In [13]:

```
la.score(x_test,y_test)
```

Out[13]:

0.9999985393411976

Elastic regression

In [16]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[16]:

ElasticNet()

In [17]:

```
print(en.intercept_)
```

-0.3961395378407815

In [18]:

```
predict=(en.predict(x_test))
```

In [19]:

```
print(en.score(x_test,y_test))
```

0.9999985249381641

Evaluation matrices

In [20]:

```
from sklearn import metrics
print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,predict))
```

Mean Absolute Error: 0.32436571871614867

In [21]:

```
print("Mean Square Error:",metrics.mean_squared_error(y_test,predict))
```

Mean Square Error: 0.16865415490153537

In [22]:

```
print("Root Mean Square Error:",np.sqrt(metrics.mean_squared_error(y_test,predict)))
```

Root Mean Square Error: 0.41067524262065686

In []: